

BALAGANESH S B

Associate Software Engineer | Data Engineer | Databricks Certified Data Engineer Associate

Avadi, Tamil Nadu, India

+91-9884549692 sbbalaganesh2004@gmail.com

linkedin.com/in/balaganesh-sb github.com/BALA1403 bala1403.github.io/portfolio

Summary

Databricks Certified Data Engineer Associate and Associate Software Engineer with hands-on experience designing scalable batch and streaming data pipelines using Databricks, PySpark, Delta Lake, and Spark SQL. Experienced in building Medallion Architecture, Auto Loader, Delta Live Tables (Lakeflow), Structured Streaming, and metadata-driven ETL frameworks. Strong foundation in modern lakehouse architecture, real-time analytics, and production-ready data engineering solutions.

Technical Skills

Programming: Python • SQL • Java

Data Engineering: Databricks • PySpark • Spark SQL • Delta Lake • Auto Loader • Delta Live Tables (Lakeflow) • Structured Streaming • Delta MERGE • Medallion Architecture • Unity Catalog • SCD Type 1/2 • Metadata-Driven ETL • Data Quality

Cloud: Microsoft Azure

Tools: Databricks Workflows • Lakeview • Git • GitHub • Jupyter Notebook

Databases: MySQL • SQL Server

Projects

Metadata-Driven ETL Framework | *PySpark, Delta Lake, Auto Loader, Unity Catalog* | [GitHub](#) **2026**

- Built a config-driven ingestion framework where a control table drives format, keys, and SCD strategy per source, so onboarding a new entity requires zero code changes.
- Implemented generic Auto Loader ingestion and hash-based SCD Type 1/2 merge logic dispatched dynamically by load type, plus a data-quality gate that quarantines rows with missing keys.
- Built run-level audit logging and pipeline health monitoring, orchestrated end-to-end using Databricks Workflows.

Enterprise Retail Lakehouse Platform | *PySpark • DLT • Auto Loader • Structured Streaming* | [GitHub](#) **2026**

- Built an end-to-end retail lakehouse (Bronze/Silver/Gold) processing customer, product, order, inventory, and clickstream data using Auto Loader and Delta Live Tables.
- Implemented declarative data-quality expectations and SCD Type 1 dimension handling via **apply_changes**, avoiding hand-written MERGE logic.
- Built a parallel Structured Streaming path for clickstream events with watermarked windowed aggregation, orchestrated across notebook and pipeline tasks via Databricks Workflows.

Banking Streaming Analytics Platform | *Structured Streaming, Delta Lake, Auto Loader* | [GitHub](#) **2026**

- Built a real-time banking transaction pipeline using Structured Streaming and Auto Loader over Unity Catalog Volumes as the streaming source.
- Implemented watermarking, stateful deduplication, and row-level plus 5-minute sliding-window fraud detection, with idempotent Delta MERGE upserts via **foreachBatch**.
- Built real-time channel KPI aggregation and fraud monitoring using tumbling windows, with live Lakeview dashboards and Databricks Workflows.

Experience

Accenture **Mar 2026 – Present**
Associate Software Engineer *Chennai, India*

- Developed ETL pipelines using Databricks, PySpark, and Spark SQL as part of enterprise data engineering training and project assignments.
- Built Medallion Architecture (Bronze–Silver–Gold) pipelines using Delta Lake, Auto Loader, and Delta Live Tables.
- Applied Agile and SDLC practices while collaborating with cross-functional teams using Git.
- Completed Enterprise Data Engineering Training (Oct 2025 – Jan 2026) in SQL, Python, PySpark/Databricks, and Azure fundamentals prior to joining.

Education

Panimalar Engineering College, Tamil Nadu **Jul 2021 – Apr 2025**
B.Tech – Computer Science and Business Systems *CGPA: 8.97 / 10.0*

Certifications

Databricks Certified Data Engineer Associate – Databricks	May 2026
Databricks Fundamentals Accreditation – Databricks Academy	Apr 2026
Databricks Generative AI Fundamentals – Databricks Academy	Apr 2026
Databricks AI Agent Fundamentals – Databricks Academy	Apr 2026
Oracle Certified Java SE 11 Developer – Oracle	Apr 2024
Oracle Certified Generative AI Professional – Oracle	Jul 2024